

---

# Fabrik: An online collaborative neural network editor

---

Anonymous Author(s)

Affiliation

Address

email

## Abstract

1 We present Fabrik, a neural network editor that provides tools to visualize, edit,  
2 and share networks in the browser. Fabrik provides an intuitive GUI to import  
3 neural networks written in popular deep learning frameworks and allows users  
4 to interact with, build, and edit models via simple drag and drop. It is designed  
5 to be framework agnostic and support high interoperability, and can be used to  
6 export models back to any supported framework. Finally, it provides powerful  
7 collaborative features for users to iterate over model design remotely and at scale.

## 8 1 Introduction

9 In recent years, long strides have been made in AI, primarily propelled by the advent of deep  
10 learning [LeCun et al., 2015]. This progress has been catalyzed by the development of well designed  
11 and documented, open source frameworks for deep learning. Over 23 frameworks have been released  
12 as of last year [KDNuggets2017], several of which have distinct strengths, such as deployment and  
13 serving [Abadi et al., 2016], inference on mobile devices [Caffe2, 2018], or extensibility for rapid  
14 research and prototyping [Paszke et al., 2017]. However, keeping up with the latest frameworks  
15 presents a daunting challenge to newcomers and experienced researchers alike. Having to constantly  
16 learn new frameworks (say, to run a model written using an unfamiliar framework), regularly proves  
17 to be a significant time sink.

18 At the same time, neural networks developed in such frameworks are often designed *collaboratively*  
19 over the course of several iterations. However, the process of this model design and development  
20 is still largely restricted to whiteboard discussions, and scaling such collaboration to remote and  
21 online settings remains an open problem. Considering the popularity of deep learning, a tool for  
22 collaborative design of neural networks is likely to have widespread utility.




Figure 1: We present Fabrik, an online collaborative neural network editor to visualize, edit, and share networks.

23 To overcome these challenges, we present Fabrik, an online collaborative neural network editor.  
24 Fabrik is an open source platform to visualize, edit, and share neural networks. See Figure 1.

25 **Visualize.** Fabrik provides a simple GUI to create networks via drag and drop or import neural  
26 networks written in popular frameworks such as Caffe, Keras, and TensorFlow.<sup>1</sup>

27 **Edit.** Fabrik supports editing layers and parameters of the neural network right in the browser. Once  
28 completed, the network can be exported to any target framework providing easy interoperability.

29 **Share.** Fabrik provides powerful collaborative features for researchers to brainstorm and design  
30 models remotely. A researcher can generate a shareable link for a model, and can see real time  
31 updates and comments from their collaborators.

32 Fabrik is targeted at newcomers and experienced deep learning practitioners alike. For newcomers,  
33 it lowers the barrier to entry to the field by enabling them to interact with neural networks online  
34 without having to grapple with syntax. For researchers, it provides a tool to easily visualize the  
35 depth and breadth of their network, verify correctness of implementation, and further allows porting  
36 between frameworks to enable them to work in their preferred framework. Finally, it provides  
37 collaborative features to brainstorm about model design remotely and at scale, and to share interactive  
38 visualizations of their models to aid in reproducibility.

## 39 2 Related Work

40 **Deep learning in the browser.** A number of browser-based deep learning tools have been released  
41 recently. Some are geared towards providing a platform to completely manage and run deep learning  
42 experiments [Aetros, 2018, KerasJS, 2017]. However, these are either paid [Aetros, 2018], or  
43 restricted to a single framework [KerasJS, 2017]. Yet others, like [ConvnetJS, 2016] have been  
44 designed primarily as teaching tools and do not plug into existing popular frameworks. In this work  
45 we propose Fabrik as a tool for easy visualization of models from a range of popular frameworks.

46 **Visualizing Neural Networks.** A few tools designed for visualizing neural networks have been  
47 proposed. For example, [Netscope, 2017] supports visualizing arbitrary networks available in Caffe’s  
48 prototxt format, however, it works only with Caffe models, and only solves a part of the problem.  
49 [Tensorboard, 2018] is a web application designed for debugging and visualizing models written in  
50 TensorFlow. While a powerful tool, it requires writing code to set up, and is thus not accessible for  
51 newcomers. Moreover, it only works with TensorFlow and related frameworks such as Keras. Fabrik  
52 supports interactively building models from within the browser, collaborating with other researchers,  
53 and exporting to popular frameworks, without having to write a line of code.

54 **Cross-framework support.** A few tools [ONNX, 2018, MMDnn, 2018] for deep learning framework  
55 interoperability have been recently proposed. [ONNX, 2018] provides an open source format for AI  
56 models, and is supported by many frameworks but it is not yet adopted by TensorFlow, currently,  
57 the most popular deep learning framework. We view ONNX as an alternative intermediate format  
58 that be integrated into our application in the future. Similarly, [MMDnn, 2018] is a cross-framework  
59 solution to convert between and visualize deep neural network models written in several frameworks.  
60 However, it does not provide tools for building models from within the browser. Additionally, it is  
61 not an online service and is not designed for remote collaboration.

## 62 3 Application Overview

### 63 3.1 Presentation Layer

64 Figure 2 shows the landing screen displayed with the GoogLeNet [Szegedy et al., 2015] network  
65 being visualized. We can see there are two distinct sections – on the right is the *Canvas* where  
66 the neural network is displayed and on the left is the *Dashboard*, which provides all of Fabrik’s  
67 core functionalities, such as adding layers, importing or exporting models, or sharing models with  
68 collaborators.

#### 69 3.1.1 Canvas

70 The canvas displays the neural network being visualized. This uses the jsPlumb [jsPlumb] library that  
71 allows for a simple interface for connecting UI elements. The layers of the network are represented  
72 as nodes of a graph which are connected based on layer dependencies. These nodes can be dragged

---

<sup>1</sup>Note that these are the top three frameworks by popularity based on Github metrics [Zacharias et al., 2018].

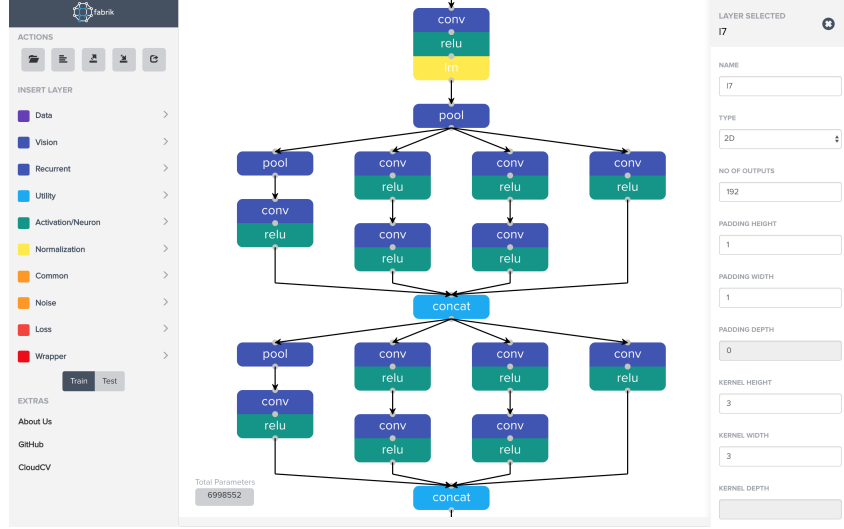


Figure 2: Visualizing the GoogLeNet [Szegedy et al., 2015] architecture with Fabrik.

and re-positioned on the canvas. Upon hovering over a layer, an overview of layer parameters is displayed as a tool-tip, and these layer parameters can be further edited by clicking on the layer and editing the values in the displayed side-pane.

### 3.1.2 Dashboard

The dashboard provides controls for all core functionalities – export, import, building, and sharing. Through the top menu in this the user can access the Model Zoo, a collection of state-of-art deep learning models that we provide inbuilt support for (Table 1 provides a complete list). Other buttons in the dashboard allow the user to initiate import / export(prompted to select source/target framework) of their neural network model and for sharing the current model with other users through a unique generated URL.

Below the top menu are layers that can be dragged and dropped onto the canvas to build or edit a model. These are collapsed into categories, each of which drills down to a set of layers. The added layers are initialized with default parameters which can be edited by the user.

### 3.1.3 Visualization Algorithm

The visualization algorithm plots the layer objects on the canvas. The main goal of the algorithm is to assign coordinates and establish connections between layer objects to yield an accurate and aesthetic visualization. To do this, the algorithm focuses on minimizing the number of overlapping layers and connections. A few challenges that the algorithm addresses include:

**Overlapping layers:** To achieve non-overlapping position allocation, we implement a hashing technique: maintaining a constant height and width for each layer object, an initial position allocation is computed. Each time a layer has been assigned co-ordinates, the immediate area surrounding the layer is searched with the help of hashed co-ordinates of visited layers to detect overlap. If overlap is detected, a constant displacement is added to the y-coordinate of the layer, and we then check for overlap recursively. The search space for detecting overlap is reduced significantly via hashing, leading to much faster visualization.

**Order of traversal:** Topological sorting in depth-first order is used for the purpose of assigning non-overlapping positions. This ensures that a layer object is visited only after all of its parent layers have been assigned positions.

To ensure non-overlapping layer connections various parameters are considered, e.g. the number of child nodes of a layer, position of the source layer etc. Based on these parameters, the slope of the connection arrow is calculated, and if needed the line is divided into sub parts to ensure that connection overlap and cutting connections are minimized.

### 105 3.2 Intermediate Representation Layer

106 The intermediate representation allows us to decouple the display portion of the application from  
107 framework specific back-ends. On importing a model from any of the supported frameworks,  
108 the layers are stored as the corresponding layers in the intermediate representation. We use the  
109 prototxt-based model configuration protocol employed by Caffe [Jia et al., 2014] as our intermediate  
110 representation as it is clean and extensible.

111 As different frameworks might have slightly varying parameters or parameter names, we first map  
112 them to corresponding names in this representation. Note that we extend the protocol to hold the  
113 union of all parameters required by all supported frameworks, so that no information is lost.

114 Creating an intermediate representation that serves all supported frameworks well is a challenging  
115 task. For example, the Caffe framework uses numerical padding, to pad the inputs before applying  
116 operations such as Convolution or Pooling. But frameworks like Keras and Tensorflow use 'SAME' or  
117 'VALID' padding, where the amount of padding required is automatically determined based on input  
118 and output sizes. To map from both types of frameworks to the same intermediate representation, we  
119 first need to convert the 'SAME'/'VALID' type of padding to its numerical counterpart by determining  
120 the input and output sizes for each layer, and perform the reverse operation when converting back.  
121 Note that while we provide a single example here, many such corner cases exist and are appropriately  
122 handled by the back-end layer.

### 123 3.3 Back-End Layer

#### 124 3.3.1 Model conversion

125 The intermediate representation also helps support conversion from one framework to another.  
126 Without such an intermediate representation, separate logic would be required to port between each  
127 pair of supported frameworks, but now the export logic is only required to be written from the  
128 intermediate representation to each of the supported framework. This reduces the problem complexity  
129 from  $O(n^2)$  to  $O(n)$ , for  $n$  different frameworks.

130 In some cases, model conversion is restricted by major differences between frameworks. For example,  
131 the *Local Response Normalization (LRN)* layer is only available in Caffe, and so it is not natively  
132 possible to port an AlexNet [Krizhevsky et al., 2012] network defined in Caffe (which uses this  
133 layer) to Keras or Tensorflow. In some cases we complete such conversions by using third-party  
134 implementations. For *LRN*, we use such implementations to allow export to these frameworks.  
135 However, there still remain cases where even this is not possible. For example, Caffe supports a  
136 *Python* layer that allows users to implement any custom logic. This can clearly not be ported to other  
137 frameworks, and in this case we report an appropriate error message.

138 Note that since TensorFlow is a supported backend of Keras, we define porting logic for just two  
139 frameworks, Caffe and Keras. For TensorFlow conversion, we first export the model as Keras and  
140 then internally use the Keras backend to port to TensorFlow.

#### 141 3.3.2 Asynchronous tasks for Model Export

142 For a model to be exported to a target framework, a back-end task need to be executed which handles  
143 the parsing of the intermediate model representation. Waiting for the task to complete might prevent  
144 a user from being able to perform any operations on the canvas. To overcome this, asynchronous  
145 tasks for model export are employed.

146 The asynchronous tasks are implemented using Celery [Celery, 2018], which allows for distributing  
147 work across threads or machines. Figure 3 shows how celery workers are used in the Fabrik  
148 infrastructure. As the call to the export event is implemented in an asynchronous manner, there is no  
149 latency at the front-end and the user can continue their operations. Once the task completes at the  
150 back-end it sends a response containing the result of the task with the help of web-sockets and the  
151 user is notified of the response on the front-end.

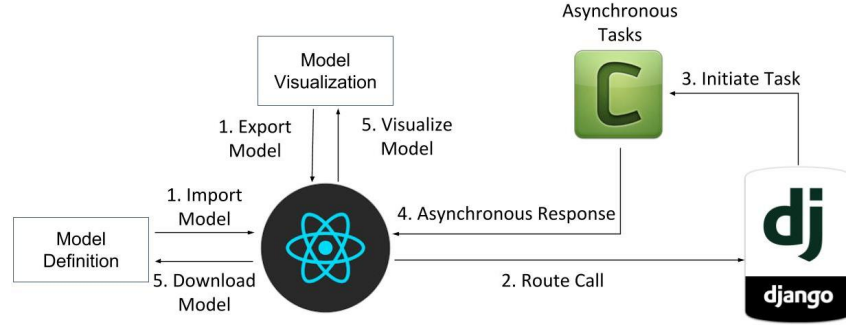


Figure 3: Export / Import Workflow

### 3.3.3 Real-time collaboration

Fabrik provides support for collaboratively building deep learning models with multiple users working in real time on a single model. For sharing a model with multiple users, a unique link for a model is generated and any user with access to the open link can view, edit, or comment on a model. A global copy of a single shared model is maintained in the database and all updates over multiple sessions are performed on this copy, using communication via web-sockets. Whenever an update operation is performed, an event message is broadcast by the back-end layer which is used by presentation layer to re-render the canvas on all active sessions. Additionally, the layer being updated is highlighted with the name of the user performing the update.

In order to achieve real-time collaboration, we employ an efficient implementation of operational transformation for the intermediate representation of a model. By storing it in JSON format, the size of each update is minimized. Based on the type of operation being performed, there are 4 types of operational update events – parameter update, layer addition, layer deletion, and layer highlight, each of which trigger an update event and the change is propagated through all sessions.

## 4 Walk-through and Use Cases

We now provide a complete walk through of the visualizing, editing, and sharing functionalities via a concrete example. Consider that we want to visualize and modify the well known VGG-16 network proposed in [Simonyan and Zisserman, 2014]. VGG is a 16-layer convolutional neural network proposed for image classification on the ImageNet dataset that is primarily composed of convolutional, pooling, and fully connected layers.

Figure 4 walks through this use case: The network is first loaded into the application through the Import option which supports a few ways: First, a user can directly upload a model configuration file as *prototxt* (for Caffe), *JSON* (for Keras), or *pbtxt* (for Tensorflow). Note that for TensorFlow and Keras that do not decouple model definition like Caffe does, serializing the model configuration into these formats is easily achieved by functions available in both frameworks. Alternatively, the model configuration can be directly pasted as text into an input form-field. Finally, model configurations available as a URL on GitHub or as a Gist can also be directly imported. As seen in Figure 4(a), the URL option is selected for import. 4(e) shows the actual URL, in this case we directly load the VGG-16 model from a prototxt file hosted on GitHub. If the import is successful, the network will be visualized on the canvas as shown in 4(b). Once the network is loaded, the user can scroll to explore different parts and zoom in/out to see a more local/global view of the network. The user can also hover over any layer to get an overview of layer parameters ( 4(b)). If on analyzing the layer parameters, the user wishes to modify a parameter, they can choose to click on that layer. This will launch a parameter pane on the right of the canvas where all layer parameters can be updated as shown in 4(c). In this case, the number of outputs for the selected convolutional layer are being modified. Please note, that as the layer was a 2D convolution, the depth dimension (Padding depth) is blocked which minimizes the chance of the user making an error.

The user can also add a new layer to an existing network by dragging it to the desired position or just clicking on it from the dashboard, the latter would attach the layer to the deepest node in the network.

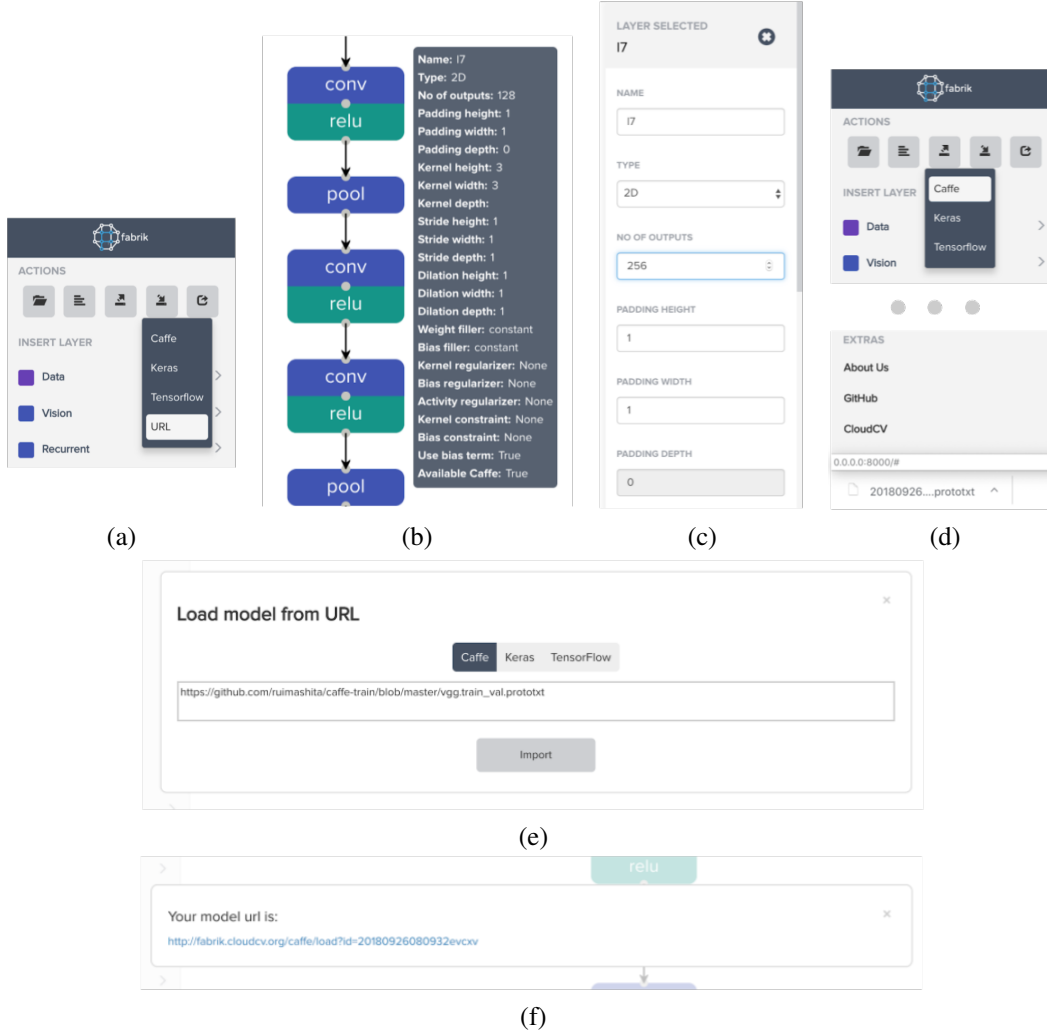


Figure 4: Screen-shots of the application depicting the stages in the life-cycle of loading(a,e), visualizing(b), editing(c), saving(d), and sharing(f) a neural network model with Fabrik.

Further, Fabrik displays the real time parameter count of the network, which updates when a layer parameter is edited or a new layer is added or deleted. This count is displayed on the bottom-left of the canvas and can provide the user an idea of the impact a layer can have on the parameter count of the network.

Upon editing the network to their satisfaction, the user can choose to export it to any of the supported frameworks or share a link to it with collaborators. Figure 4(d) shows the network being exported back to the Caffe framework from which it was originally loaded. At the bottom-left of the screen we can see the exported output being downloaded as a prototxt file. Lastly, figure 4 (f) shows the unique link generated for this model, that can be used to access it in the future, or be shares with others.

Instead of exporting the network to the Caffe framework from which it was imported, the user could also export it to a different supported framework, such as Keras or Tensorflow. Table 1 shows the results of model conversion between currently supported frameworks. As seen, certain models cannot be converted to certain frameworks. For example, exporting GoogLeNet to TensorFlow and Keras is not possible due to lack of support for the LRN layer. Lastly, the user can also choose the share the model to a link, which can be used simultaneously by multiple users to collaboratively view and edit the network and see real-time updates.

Models	Caffe	Keras	TensorFlow
Recognition			
Inception V3 [Szegedy et al., 2016]	✓	✓	✓
Inception V4 [Szegedy et al., 2017]	✓	✓	✓
ResNet 101 [He et al., 2016]	✓	✓	✓
VGG 16 [Simonyan and Zisserman, 2014]	✓	✓	✓
GoogLeNet [Szegedy et al., 2015]	✓	✗	✗
SqueezeNet [Iandola et al., 2016]	✓	✗	✗
AllCNN [Springenberg et al., 2014]	✓	✗	✗
DenseNet [Huang et al., 2017]	✓	✗	✗
AlexNet [Krizhevsky et al., 2012]	✓	✓	✓
Detection			
YoloNet [Redmon et al., 2016]	✓	✓	✓
FCN32 Pascal [Long et al., 2015]	✓	✗	✗
Seq2Seq			
Pix2Pix [Isola et al., 2017]	✓	✗	✗
Visual Question Answering			
VQA [Antol et al., 2015]	✓	✓	✓

Table 1: List of tested models on Fabrik

## 5 Conclusion and Future Work

In this work we propose Fabrik, a simple GUI for visualizing, creating, editing, and sharing neural network models in the browser. Fabrik provides novel real-time collaboration features that to assist researchers to brainstorming and design models remotely and at scale.

It provides a framework agnostic visualization scheme based on a novel intermediate representation that can be used to quickly and efficiently visualize model architectures. While of today it supports three popular deep learning frameworks, as future work we plan on adding support for additional popular frameworks, such as PyTorch and Caffe2. Further, we plan to incorporate ONNX as our intermediate representation as a step towards long term interoperability and maintainability. Finally, Fabrik will be available as an online service which provides instant access on any browser supporting device for the purposes of debugging, reproducibility, learning, and collaboration.

## References

- Martín Abadi, Paul Barham, Jianmin Chen, Zhifeng Chen, Andy Davis, Jeffrey Dean, Matthieu Devin, Sanjay Ghemawat, Geoffrey Irving, Michael Isard, et al. Tensorflow: a system for large-scale machine learning. In *OSDI*, volume 16, pages 265–283, 2016.
- Aetros. Aetros: Deep learning experiment management platform. <https://aetros.com/>, 2018.
- Stanislaw Antol, Aishwarya Agrawal, Jiasen Lu, Margaret Mitchell, Dhruv Batra, C Lawrence Zitnick, and Devi Parikh. Vqa: Visual question answering. In *Proceedings of the IEEE international conference on computer vision*, pages 2425–2433, 2015.
- Caffe2. Caffe2: A new lightweight, modular, and scalable deep learning framework. <https://caffe2.ai/>, 2018.
- Celery. Celery: Distributed task queue. <https://github.com/celery/django-celery>, 2018.
- ConvnetJS. ConvnetJS: Deep learning in javascript. <https://github.com/karpathy/convnetjs>, 2016.
- Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 770–778, 2016.

234 Gao Huang, Zhuang Liu, Laurens Van Der Maaten, and Kilian Q Weinberger. Densely connected  
235 convolutional networks. In *CVPR*, volume 1, page 3, 2017.

236 Forrest N Iandola, Song Han, Matthew W Moskewicz, Khalid Ashraf, William J Dally, and Kurt  
237 Keutzer. Squeezenet: Alexnet-level accuracy with 50x fewer parameters and < 0.5 mb model size.  
238 *arXiv preprint arXiv:1602.07360*, 2016.

239 Phillip Isola, Jun-Yan Zhu, Tinghui Zhou, and Alexei A Efros. Image-to-image translation with  
240 conditional adversarial networks. *arXiv preprint*, 2017.

241 Yangqing Jia, Evan Shelhamer, Jeff Donahue, Sergey Karayev, Jonathan Long, Ross Girshick, Sergio  
242 Guadarrama, and Trevor Darrell. Caffe: Convolutional architecture for fast feature embedding.  
243 In *Proceedings of the 22nd ACM international conference on Multimedia*, pages 675–678. ACM,  
244 2014.

245 jsPlumb. jsPlumb: Visual connectivity for webapps. <https://github.com/jsplumb/jsplumb>.

246 KDNuggets2017. Ranking popular deep learning frameworks  
247 for data science. [https://www.kdnuggets.com/2017/10/  
248 ranking-popular-deep-learning-libraries-data-science.html](https://www.kdnuggets.com/2017/10/ranking-popular-deep-learning-libraries-data-science.html).

249 KerasJS. KerasJS: Run keras models in the browser. [https://github.com/transcranial/  
250 keras-js](https://github.com/transcranial/keras-js), 2017.

251 Alex Krizhevsky, Ilya Sutskever, and Geoffrey E Hinton. Imagenet classification with deep convolu-  
252 tional neural networks. In *Advances in neural information processing systems*, pages 1097–1105,  
253 2012.

254 Yann LeCun, Yoshua Bengio, and Geoffrey Hinton. Deep learning. *nature*, 521(7553):436, 2015.

255 Jonathan Long, Evan Shelhamer, and Trevor Darrell. Fully convolutional networks for semantic  
256 segmentation. In *Proceedings of the IEEE conference on computer vision and pattern recognition*,  
257 pages 3431–3440, 2015.

258 MMDnn. MMDnn: A comprehensive, cross-framework solution to convert, visualize and diagnosis  
259 deep neural network models. <https://github.com/Microsoft/MMdnn>, 2018.

260 Netscope. Netscope: A web-based tool for visualizing neural network topologies. [https://github.  
261 com/ethereon/netscope](https://github.com/ethereon/netscope), 2017.

262 ONNX. ONNX: Open neural network exchange format. <https://onnx.ai/>, 2018.

263 Adam Paszke, Sam Gross, Soumith Chintala, and Gregory Chanan. Pytorch, 2017.

264 Joseph Redmon, Santosh Divvala, Ross Girshick, and Ali Farhadi. You only look once: Unified,  
265 real-time object detection. In *Proceedings of the IEEE conference on computer vision and pattern  
266 recognition*, pages 779–788, 2016.

267 Karen Simonyan and Andrew Zisserman. Very deep convolutional networks for large-scale image  
268 recognition. *arXiv preprint arXiv:1409.1556*, 2014.

269 Jost Tobias Springenberg, Alexey Dosovitskiy, Thomas Brox, and Martin Riedmiller. Striving for  
270 simplicity: The all convolutional net. *arXiv preprint arXiv:1412.6806*, 2014.

271 Christian Szegedy, Wei Liu, Yangqing Jia, Pierre Sermanet, Scott Reed, Dragomir Anguelov, Du-  
272 mitru Erhan, Vincent Vanhoucke, and Andrew Rabinovich. Going deeper with convolutions. In  
273 *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 1–9, 2015.

274 Christian Szegedy, Vincent Vanhoucke, Sergey Ioffe, Jon Shlens, and Zbigniew Wojna. Rethinking  
275 the inception architecture for computer vision. In *Proceedings of the IEEE conference on computer  
276 vision and pattern recognition*, pages 2818–2826, 2016.

277 Christian Szegedy, Sergey Ioffe, Vincent Vanhoucke, and Alexander A Alemi. Inception-v4, inception-  
278 resnet and the impact of residual connections on learning. In *AAAI*, volume 4, page 12, 2017.



- 279 Tensorboard. Tensorboard: Suite of web applications for inspecting and understanding your tensorflow  
280 runs and graphs. <https://github.com/tensorflow/tensorboard>, 2018.
- 281 Jan Zacharias, Michael Barz, and Daniel Sonntag. A survey on deep learning toolkits and libraries  
282 for intelligent user interfaces. *arXiv preprint arXiv:1803.04818*, 2018.